

Qoneqt Video Factory

Topic in. Global Feed short out.

Theme: Build an LLM-Powered Content Pipeline for Qoneqt. A repeatable system that turns a topic into a publish-ready, captioned vertical video for the Qoneqt Global Feed.

TEAM

[Team Name] · [Team ID]

COLLEGE

[College Name]

CORTISOL BY FIFTEEN PERCENT,

CHAOS ERUPTS? THEY

COFFEE KA CAFFEINE

Three shorts generated by the pipeline tonight:
Tech, Motivation, Hinglish presets

The Global Feed is hungry. Video is slow.

What is the problem?

A social feed lives on fresh video every day. Making one 30-second short by hand means writing a script, finding visuals, recording a voice, editing, and adding captions. A creator manages one a day. The Global Feed needs hundreds.

2-4 hrs

to hand-make one 30-second short

0

Qoneqt-native topic-to-video tools

Result: thin, repetitive video content, creators who burn out, and communities that go quiet between posts.

?

Who is affected

Qoneqt creators and community admins who must post daily, and the platform, whose Global Feed runs on steady video.

~

Current solutions

Manual editing apps (CapCut, Canva), one-shot AI video generators on paid credits, or hiring an agency. All produce one video at a time.

!

Key gaps

- Not repeatable: no batch, no pipeline
- Not community-aware: generic tone
- Paid per video; no enforced hook or captions

A video factory, not a video button

Type a topic. Pick a community. Get a publish-ready short and its post text in about a minute.

1 Plan

The LLM writes 3 hooks, scores them, keeps the best, then a 5–7 scene script with a stock-search query per scene, plus caption and hashtags.

2 Produce

An AI voice reads each scene. An AI still is generated per scene, or a photo fetched, then given motion. Whisper gives word timings for word-pop captions.

3 Publish

ffmpeg assembles a 1080×1920 MP4 with burned-in captions. The app hands you the video, caption, and hashtags ready to paste into Qoneqt.

What makes it better than existing options

	Manual editing	One-shot AI tools	Video Factory
Repeatable at scale	No	No	Yes, batch of 10 per click
Community-aware tone	By hand	No	6 presets, one dict entry each
Hook, captions, post text	By hand	Partly	Built in, every video
Cost per video	Hours	Paid credits	₹0 on free tiers

Input → Process → Output



Every stage has a fallback chain, so one flaky free API never kills a job. A single worker thread drains the job queue; batch mode is just N queued topics.

Frontend

One static HTML page. Textarea for topics (one per line), community dropdown, live stage progress, inline player, Copy post text, Download.

Backend

FastAPI. POST /generate, GET /jobs, static /out. In-memory job dict, one worker thread, files per job in out/<id>/.

AI layer

LLM for plan JSON (validated, guided retry). TTS per scene. Photo or clip search per scene. Whisper word timings aligned to script words.

Data + deploy

Outputs as files (mp4, jpg, json). Docker image with ffmpeg and Noto fonts. Deployed on a free Hugging Face Space, keys as secrets.

Stack Python 3.11 · FastAPI · httpx · ffmpeg · Docker · Groq (gpt-oss-120b, Whisper large-v3-turbo) · Pollinations FLUX · Hugging Face Spaces · ElevenLabs · edge-tts · Wikimedia Commons

Five things that make it work

1 Community presets

Six Qoneqt communities (General, Tech, Fitness, Motivation, Finance, Hinglish Fun). Each sets tone, language, narrator voice, caption style, hashtags, and caption colour. Adding one is a single dict entry.

2 Self-scored hooks

The LLM writes three opening lines, scores each for scroll-stopping power, and the best one opens the video. Quality control inside the prompt, zero extra calls.

3 Word-pop captions

Whisper word timestamps drive karaoke-style captions: each spoken word lights up in the preset colour, aligned back to the script's own spelling. The look that performs on shorts.

4 Batch and fallbacks

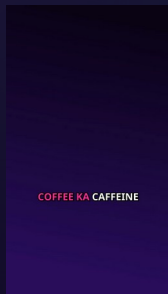
Ten topics per click, one worker, no babysitting. Every stage has a fallback (LLM, voice, stock, captions), so the factory keeps producing when a free API blinks.

5 Post-ready output

Not just an MP4. Hook, caption, and hashtags come out together, with a one-click Copy post text button, so publishing to Qoneqt is a 30-second manual step.

Real output frame

Hinglish preset, edge-tts fallback voice, Whisper-timed captions, pink accent word. Generated, not mocked.



Impact and feasibility

₹0

API cost per video on free tiers

~1 min

to make a 30–45 s short
(measured tonight)

10

topics per click in batch mode

6

community presets, one dict entry to add more

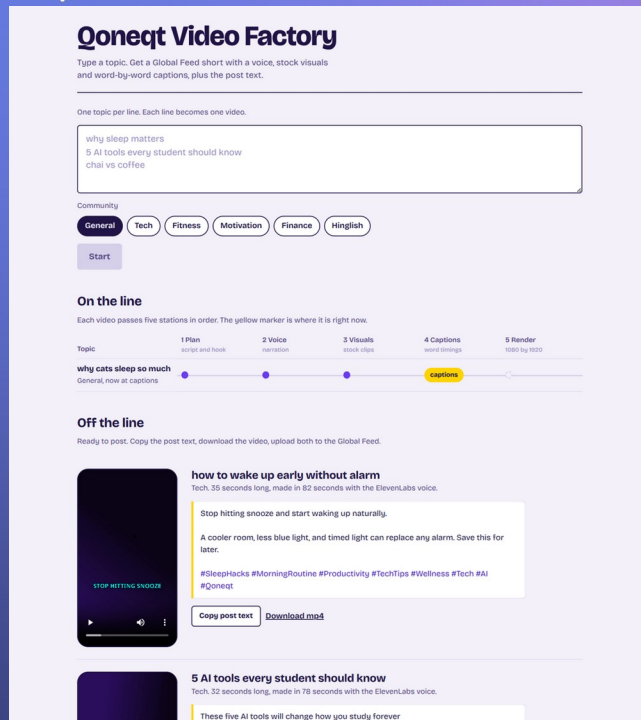
Who benefits and how

- Creators: turn ideas into daily shorts without editing skills or hours of work
- Community admins: on-tone videos in their own voice, Hinglish included
- Qoneqt: a steadier, higher-quality stream on the Global Feed, and a tool it can offer creators natively
- Real-world: 100 topics in overnight, 100 publish-ready videos out

Scalability and feasibility

- One free CPU: ~40–60 videos/hour per worker; scale by adding workers
- Free tiers: Groq LLM + Whisper, ElevenLabs (edge-tts fallback), Wikimedia visuals
- One Docker container on a free Hugging Face Space; keys stay secret
- Working end to end today; live URL and Global Feed post by the finale

Built and running tonight



What the prototype does today

- Batch of three topics queued from the page; each ran plan → images → voice → visuals → captions → render with live progress
- 34-second videos produced in about 55 seconds each, ElevenLabs voice, word-pop captions, post text ready to copy
- Hinglish preset verified on the fallback voice; every stage's fallback exercised

GitHub github.com/<user>/qoneqt-video-factory
 Live demo
huggingface.co/spaces/<user>/qoneqt-video-factory

Qoneqt doesn't need one AI video. It needs a factory. We built the factory.

Why select this: it already ships, costs nothing to run, speaks each community's language, and is built on Qoneqt's real content format, not a concept.